

EMBED: test exams in the shared latent space. Six view-availability groups, each through its own encoder over frozen ViT features.  $W_2$  values are scale-normalised. Axes are scaled per panel, because the anchors shrink the latent space by an order of magnitude or more.

