


EMBED: test exams in the shared latent space. Six view-availability groups, each through its own encoder over frozen ViT features. W_2 values are scale-normalised. Axes are scaled per panel.

No anchors

W_2 between groups 1.03 · between classes 0.64 (this seed)



Class anchors, anchor learning rate 10x (5e-4)

W_2 between groups 0.24 · between classes 0.86 (this seed)
cloud → learnt anchor 0.47 (this seed)

