

NHANES: where each test patient lands in the shared latent space. Three groups with 10, 13 and 20 features, each through its own encoder. W_2 values are scale-normalised. Axes are scaled per panel, because the anchors shrink the latent space by an order of magnitude or more.

